

Statistical Benchmark Report

Generated: 2025-11-17 19:54

This report provides quantitative evaluation results for NDCG@k and faithfulness, correlation diagnostics, outlier detection, and statistical visualizations.

1. Summary Statistics

files	100.0000
ndcg@k mean	1.0000
ndcg@k ci95 lo	1.0000
ndcg@k ci95 hi	1.0000
faith mean	0.6400
faith ci95 lo	0.6025
faith ci95 hi	0.6764
ndcg@k median	1.0000
faith median	0.6860
ndcg@k std	0.0000
faith std	0.1834
bootstrap iters	2000.0000
iqr k	1.5000
z thresh	3.0000
faith high thr	0.5000
faith mid thr	0.2500

2. Correlation and Outlier Analysis

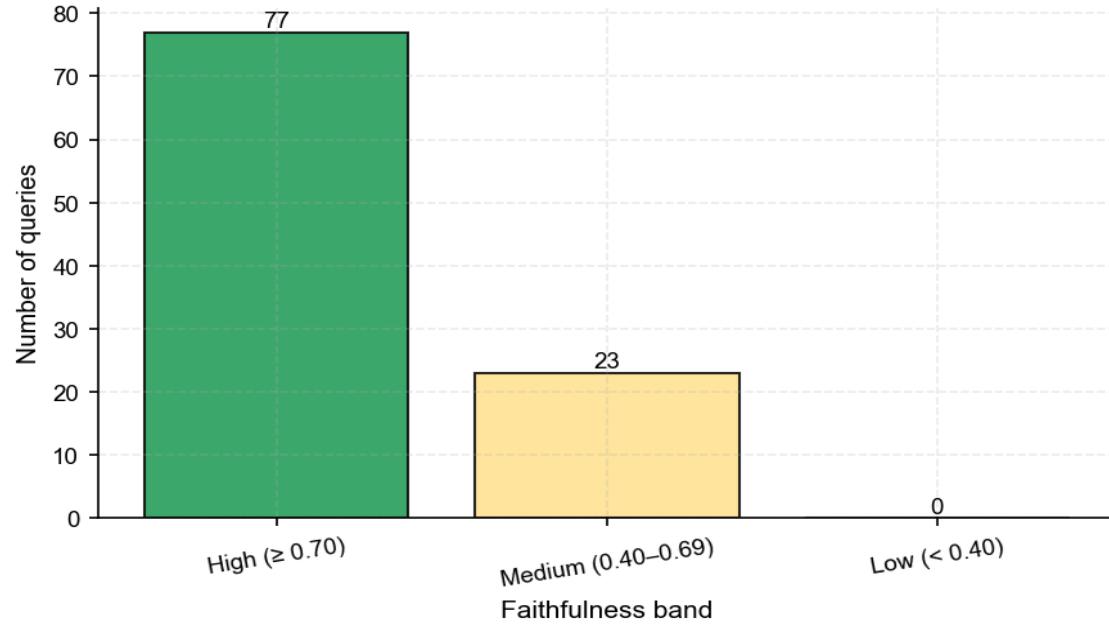
Metric Pair	Correlation	p-value
Pearson r	nan	nan
Spearman ρ	nan	nan

No outliers detected.

3. Visual Analytics

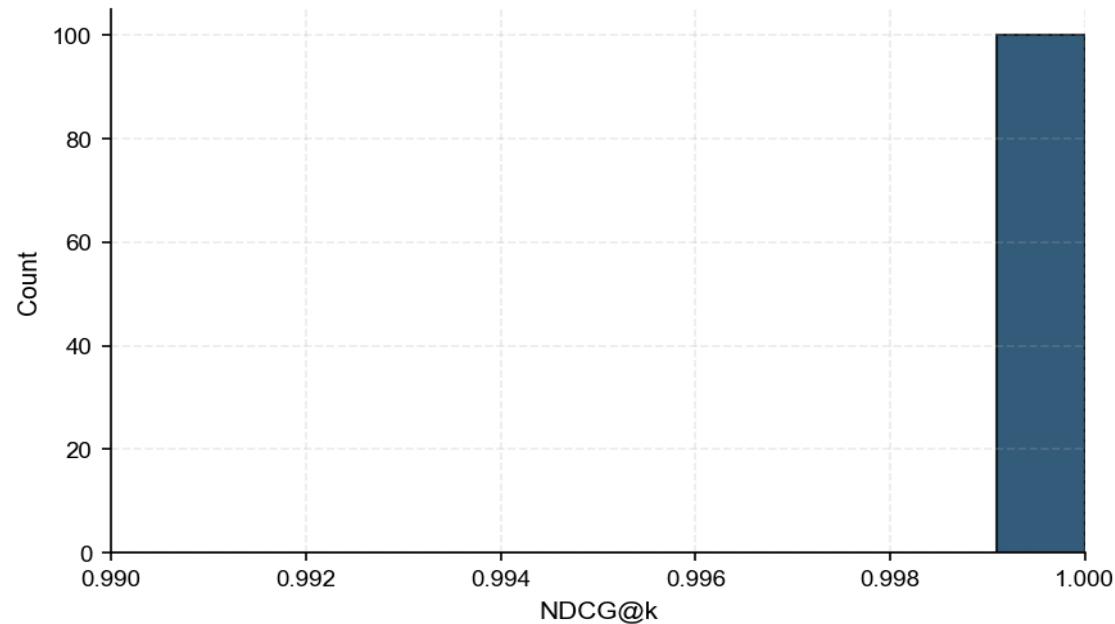
faithfulness bands global.png

Figure 2: Faithfulness bands across all queries



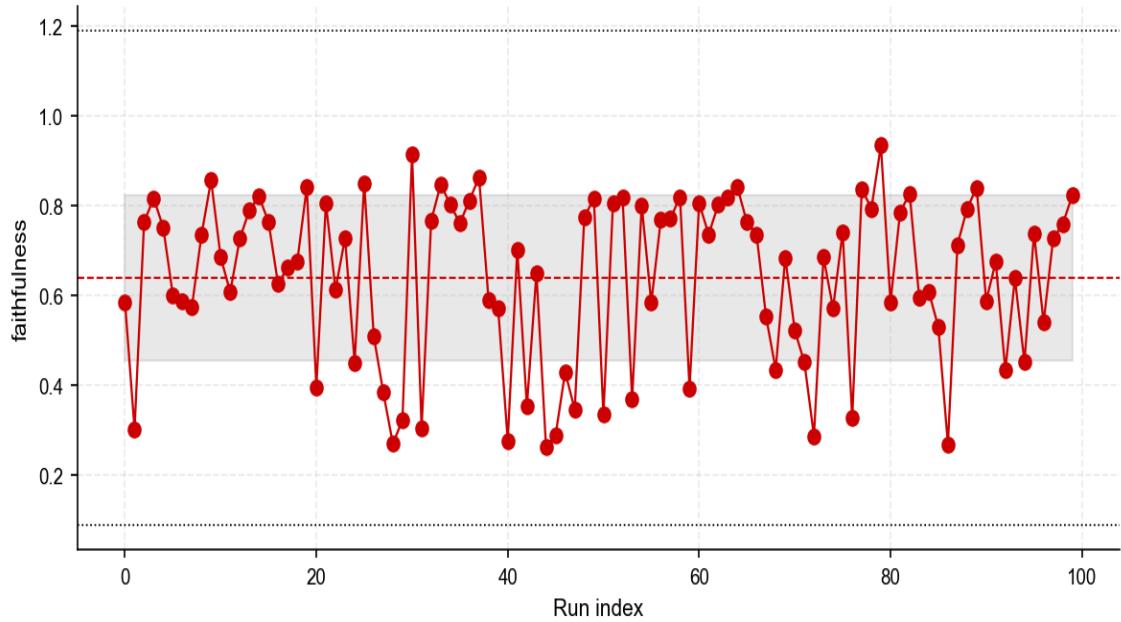
hist ndcg.png

Figure 1: Distribution of NDCG@k



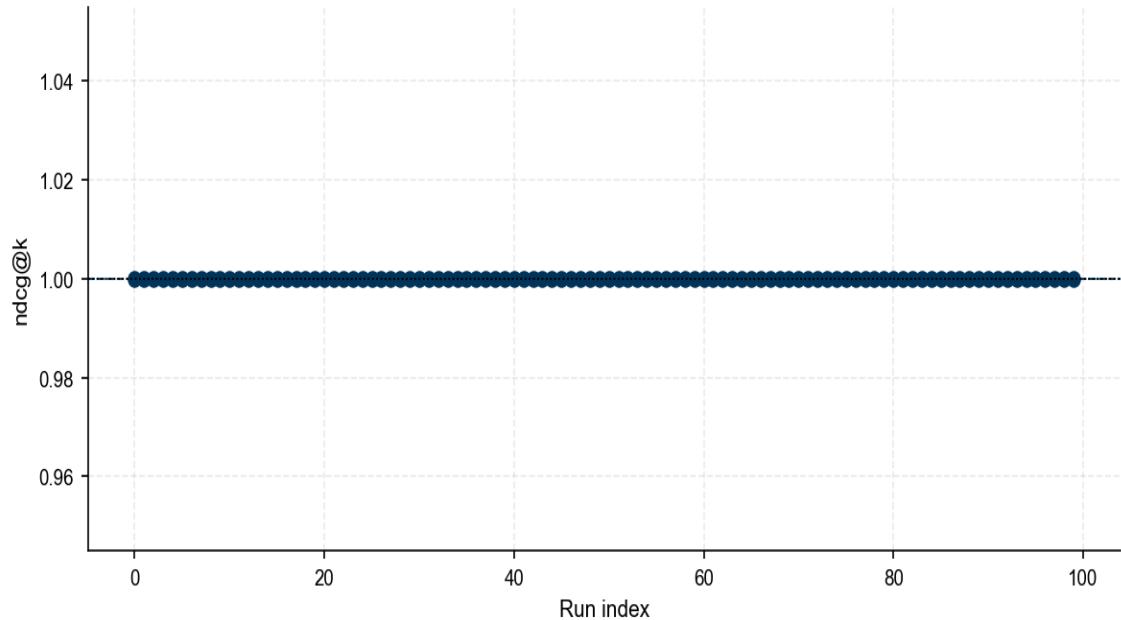
run order faithfulness.png

Figure 5: Run-order chart for faithfulness



run order ndcg@k.png

Figure 4: Run-order chart for ndcg@k



scatter ndcg vs faithfulness.png

Figure 3: Scatter NDCG@k vs Faithfulness ($r=\text{nan}$, $\rho=\text{nan}$)

